

Over-optimisation frontier: KL vs held-out reward

